Gaussian Process Regression for Gaussian Random Fields in Cosmology

Verena Alton

2024S 250155-1 Cosmic Structures: Theory, Numerics, and Statistics

July 28, 2024

1 Introduction to Gaussian Process Regression

Gaussian Process Regression (GPR) is a non-parametric, Bayesian approach to regression that is particularly useful for modelling and predicting the behaviour of complex systems where the underlying relationships between variables are unknown or highly nonlinear. GPR is based on the concept of a Gaussian process, which is a generalisation of the Gaussian distribution to infinite-dimensional spaces. [1]

1.1 Gaussian Processes

A Gaussian process (GP) is a collection of random variables which have a joint Gaussian distribution. Formally, a Gaussian process is defined as follows.

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

where $\mathbf{x}$ represents input variables, $m(\mathbf{x})$ is the mean function, and $k(\mathbf{x}, \mathbf{x}')$ is the covariance (or kernel) function. [2, 1]

1.2 Gaussian Process Regression

In GPR, we aim to make predictions about the function f given a set of observed data points $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$. The observations are assumed to be related to the latent function $f(\mathbf{x})$ via:

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

where $\epsilon_i \sim \mathcal{N}(0, \sigma_n^2)$ represents Gaussian noise with variance σ_n^2.

Given the observed data, we seek to determine the posterior distribution of the function values at new input points $\mathbf{x}^*$. The joint distribution of the observed outputs $\mathbf{y}$ and the function values $\mathbf{f}^*$ at the new inputs is given by:

$$\begin{pmatrix} \mathbf{y} \\ \mathbf{f}^* \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \mathbf{m}(\mathbf{X}) \\ \mathbf{m}(\mathbf{X}^*) \end{pmatrix}, \begin{pmatrix} K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 I & K(\mathbf{X}, \mathbf{X}^*) \\ K(\mathbf{X}^*, \mathbf{X}) & K(\mathbf{X}^*, \mathbf{X}^*) \end{pmatrix} \right)$$

where: - $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_n]^T$ is the matrix of training inputs, - $\mathbf{X}^* = [\mathbf{x}_1^*, \mathbf{x}_2^*, \ldots, \mathbf{x}_m^*]^T$ is the matrix of test inputs, - $\mathbf{m}(\mathbf{X})$ and $\mathbf{m}(\mathbf{X}^*)$ are the mean vectors of the training and test inputs, - $K(\mathbf{X}, \mathbf{X})$ is the covariance matrix for the training inputs, - $K(\mathbf{X}, \mathbf{X}^*)$ is the covariance matrix between the training and test inputs, - $K(\mathbf{X}^*, \mathbf{X}) = K(\mathbf{X}, \mathbf{X}^*)^T$, - $K(\mathbf{X}^*, \mathbf{X}^*)$ is the covariance matrix for the test inputs. [2, 1]

1.3 Posterior Distribution

To make predictions, we need the conditional distribution of $\mathbf{f}^*$ given $\mathbf{y}$. Using properties of the multivariate normal distribution, this conditional distribution is given by:

$$\mathbf{f}^* | \mathbf{X}, \mathbf{y}, \mathbf{X}^* \sim \mathcal{N}(\mu^*, \mathbf{\Sigma}^*)$$

where:

$$\mu^* = \mathbf{m}(\mathbf{X}^*) + K(\mathbf{X}^*, \mathbf{X})[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 I]^{-1}(\mathbf{y} - \mathbf{m}(\mathbf{X}))$$
$$\mathbf{\Sigma}^* = K(\mathbf{X}^*, \mathbf{X}^*) - K(\mathbf{X}^*, \mathbf{X})[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 I]^{-1} K(\mathbf{X}, \mathbf{X}^*)$$

Here, μ^* is the mean of the predictive distribution and $\mathbf{\Sigma}^*$ is the covariance of the predictive distribution. [2, 1]

2 Application of GPR to sample a Gaussian Random Field given a certain Covariance and Observations

In the lecture, we already learnt how to sample a Gaussian Random Field (GRF) with correlated noise. For this task, we used the Fourier transformation, since in Fourier space the covariance between modes is already diagonal (for a stationary field) and can therefore be expressed by the power spectrum. For a discrete set of distinct wave numbers k_i, where we defined $k_j' := -k_j$, we rewrote the definition of the power spectrum as the diagonal covariance matrix:

$$\tilde{C}_{ij} = \mathbb{E}[\tilde{\delta}(k_i)\tilde{\delta}^*(k_j)] = (2\pi)^3 P(k_i)\delta_{ij}.$$

With that, we created Gaussian random fields with correlated noise, where the degree of correlation depends on the parameter α of the power spectrum $P(k) = k^{-\alpha}$. See Figure 1 for an example of a correlated field with a power spectrum of $P(K) = k^{-3}$. [3]

Now imagine a scenario where the covariance of a GRF is known as well as the values of some observations. How could we sample a full GRF that satisfies the given covariance and matches the given observation values? Here we can use Gaussian process regression. To see the implementation of this project, please feel free to visit its Github repository.

2.1 Sampling Observations

First, we have to generate some observation values. Of course in real life applications these values come from real, empiric observations. However, in this project we will generate a GRF with Fourier transformation as explained above which will serve as our (latent) function f and then use some randomly chosen observations from it to fit our new GRF with the same (or similar) covariance as this original field and matching observations.

Figure 1: Gaussian random field with correlated noise, characterised by power spectrum $P(k) = k^{-\alpha}$ with $\alpha = 3$.

2.2 Expressing the Power Spectrum as a Covariance Kernel

Now we have to express our power spectrum as a covariance matrix which the GPR will use to generate suitable distributions. We have to convert the power spectrum, which is in Fourier space, to a covariance matrix in physical space. Therefore, we will use an inverse Fourier transformation on the power spectrum. As suggested in [4], we use that

$$\frac{1}{r^\alpha} = \frac{2\pi^{\alpha/2}}{\Gamma(\alpha/2)} \int_0^\infty \lambda^{\alpha-1} e^{-\pi\lambda^2 r^2} d\lambda$$

and derive with $\mu = 1/\lambda$

$$\frac{2\pi^{\alpha/2}}{\Gamma(\alpha/2)} \int_0^\infty \lambda^{\alpha-1} e^{-\pi|k|^2/\lambda^2} d\lambda = \frac{2\pi^{\alpha/2}}{\Gamma(\alpha/2)} \int_0^\infty \mu^{(n-\alpha)-1} e^{-\pi|k|^2\mu^2} d\mu$$

$$= \frac{2\pi^{\alpha/2}}{\Gamma(\alpha/2)} \frac{\Gamma((n-\alpha)/2)}{2\pi^{n/2-\alpha/2}} \frac{1}{|k|^{n-\alpha}}$$

$$= \frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)} \frac{1}{|k|^{n-\alpha}},$$

$$= \frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)} |k|^{\alpha-n},$$

which holds for $0 < \alpha < n$.

As we want to utilise Scikit Learn's framework for the Gaussian process regression [5], we have to express the covariance as a kernel.

2.2.1 Implementing a custom Kernel in the Scikit Learn Framework

Attempts were made to implement a custom kernel that inherits from the Scikit Learn kernel class to achieve exact results. There were several attempts to implement our desired covariance as a kernel, but all failed for various and different reasons. For example, the resulting covariance matrix was not positive definite, the Cholesky-factorisation was not possible, the optimisation of hyper-parameters did not converge or there was just no correlation resulting from the covariance kernel. There were also more ominous errors as type or shape mismatches, where the reasons/solutions probably lie deep in the code of Scikit Learn's framework. One of the more simple attempts is presented in the following. The nugget enhances numerical stability. The constant_value is the $\frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)}$ term, which is passed as an argument to the kernel for clearness reasons. It was possible to fit the GPR with it, but the prediction consisted of NaNs.

```
class PowerSpectrumKernel3D(GenericKernelMixin, Kernel):
    def __init__(self, alpha=2, constant_value=1.0, nugget=1e-5):
        self.alpha = alpha
        self.constant_value = constant_value
        self.nugget = nugget

    def __call__(self, X, Y=None, eval_gradient=False):
        if Y is None:
            Y = X

        # Calculate pairwise Euclidean distances in 3D
        dists = np.sqrt(((X[:,np.newaxis]-Y[np.newaxis,:])**2).sum(axis=2))

        # Avoid zero distances to prevent division by zero or log(0)
        dists[dists == 0] = 1e-10

        # Compute the covariance matrix
        K = self.constant_value * np.abs(dists) ** (self.alpha - 3)

        # Ensure the matrix is positive semi-definite
        K = (K + K.T) / 2
        K += np.eye(K.shape[0]) * self.nugget

        if eval_gradient:
            # Gradient with respect to the hyperparameters, if needed
            K_gradient = np.zeros((X.shape[0], Y.shape[0], 2))

            # gradient with respect to constant_value
            K_gradient[..., 0] = K / self.constant_value

            # gradient with respect to alpha
            K_gradient[..., 1] = K * np.log(np.abs(dists))
```

```
            return K, K_gradient

        return K

    def diag(self, X):
        return np.full(X.shape[0], self.constant_value + self.nugget)

    def is_stationary(self):
        return True
```

To see more attempts, please visit the Github repository of this project.

2.2.2 Using the predefined Rational Quadratic Kernel with adjusted Parameters

Due to time and capacity reasons, a more dirty approach was pursued. None of the ready-to-use kernels provided by Scikit Learn presents the desired power spectrum covariance exactly. However, the RationalQuadratic kernel has some similarity: $k_{\mathrm{RQ}}(x_i, x_j) = (1 + \frac{k(x_i,x_j)^2}{2\beta l^2})^{-\beta}$. Therefore, the parameters β and l, which can be passed to the kernel, will be defined so that it resembles our desired power spectrum. With $\beta = (n - \alpha)/2$ and $l = \frac{1}{n-\alpha}|\frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)}|^{1/(\alpha-n)}$ we get

$$k_{\mathrm{RQ}}(x_i, x_j) = (1 + \frac{k(x_i,x_j)^2}{2\beta l^2})^{-\beta}$$

$$= (1 + \frac{|k|^2}{2\frac{n-\alpha}{2}\frac{1}{n-\alpha}(|\frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)}|^{\frac{1}{n-\alpha}})^2})^{-\frac{n-\alpha}{2}}.$$

Now, neglecting the 1 in the front (which is admittedly not very rigorous), we get

$$\approx \frac{|k|^{2(-\frac{n-\alpha}{2})}}{2\frac{n-\alpha}{2}\frac{1}{n-\alpha}|\frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)}|^{\frac{1}{n-\alpha}2(-\frac{n-\alpha}{2})}}$$

$$= \frac{|k|^{2(-\frac{n-\alpha}{2})}}{|\frac{\pi^{\alpha-n/2}\Gamma((n-\alpha)/2)}{\Gamma(\alpha/2)}|^{-1}}$$

$$= |\frac{\pi^{\alpha-n/2}\Gamma((n - \alpha)/2)}{\Gamma(\alpha/2)}||k|^{\alpha-n}.$$

Of course, we will not match the desired covariance of the original field exactly, but we have an approximation.

2.3 Results

With an implementation using the RQ kernel as described above, we get the following results. We see the original field (latent function) on the top, the observations in the middle, and the Gaussian random field generated by the fitted Gaussian process regressor on the bottom.

Since the plots are only 2D, the presented results are only from one z-value (here we have $z = 16/2 = 8$).

We now want to vary the α parameters as well as the number of observations and compare the results.

2.3.1   Varying α

First, we will vary the parameter α that defines the strength of the correlations. Higher α induces stronger correlations. We have $\alpha = 3.5$ in Figure 3, $\alpha = 5.5$ in Figure 4, and $\alpha = 7.5$ in Figure 5.

We observe that our regression works better with a higher α. Probably, our approximated kernel induces stronger correlations than the original power spectrum covariance.

2.3.2   Varying the number of observations

Now, we will vary the number of observations with a fixed $\alpha = 6.5$. We have 20 observations in Figure 6, 500 observations in Figure 7, and 1000 observations in Figure 8.

As expected, more observations lead to more similar random fields. Of course in real life applications it would be difficult to gather that many observations, as in the last example we have already 100 observations which is roughly a quarter of the total grid points.

3   Conclusion

Gaussian process regression appears as a valuable tool to (re)construct a Gaussian random field given a certain covariance and observations. In this project, only an approximation to the original covariance is achieved, which could definitely be improved. With this approximation, the reconstructed GRF has stronger correlations than the original. It could be interesting to use other Machine Learning Frameworks with better fitting kernel options or even implement the Gaussian process regression or at least the covariance kernel from scratch. Furthermore, the resolution of the GRF was very low due to limited RAM, as the GPR allocated arrays of size $N * N * N$, with N being the number of grid points in one dimension. With $N = 32$, the laptop that was used to run this project crashed. The implementation of this project can be found in Github. [6]
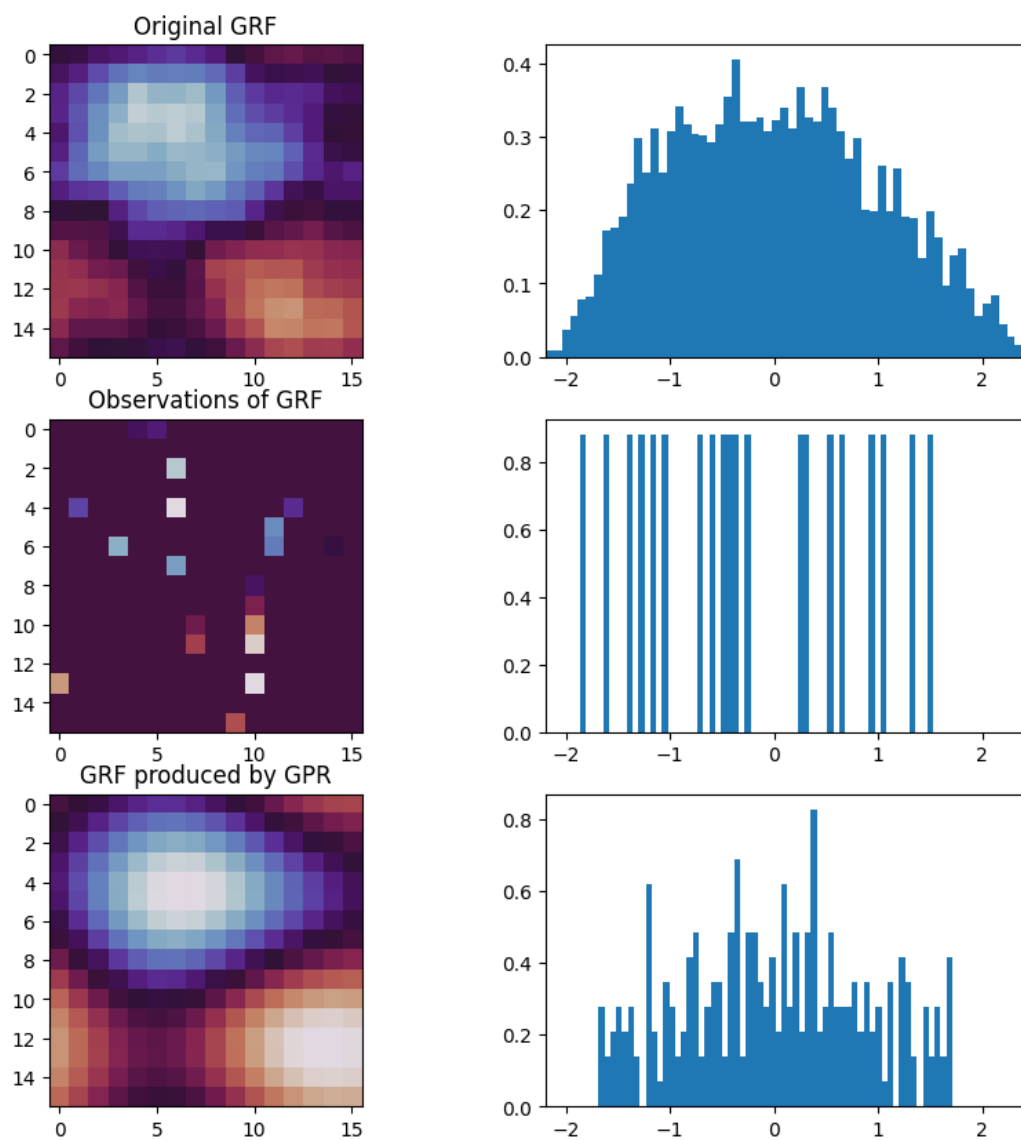
Figure 2: Gaussian process regression with $\alpha = 5.5$ and 300 observations.

Figure 3: Gaussian process regression with $\alpha = 3.5$ and 300 observations.
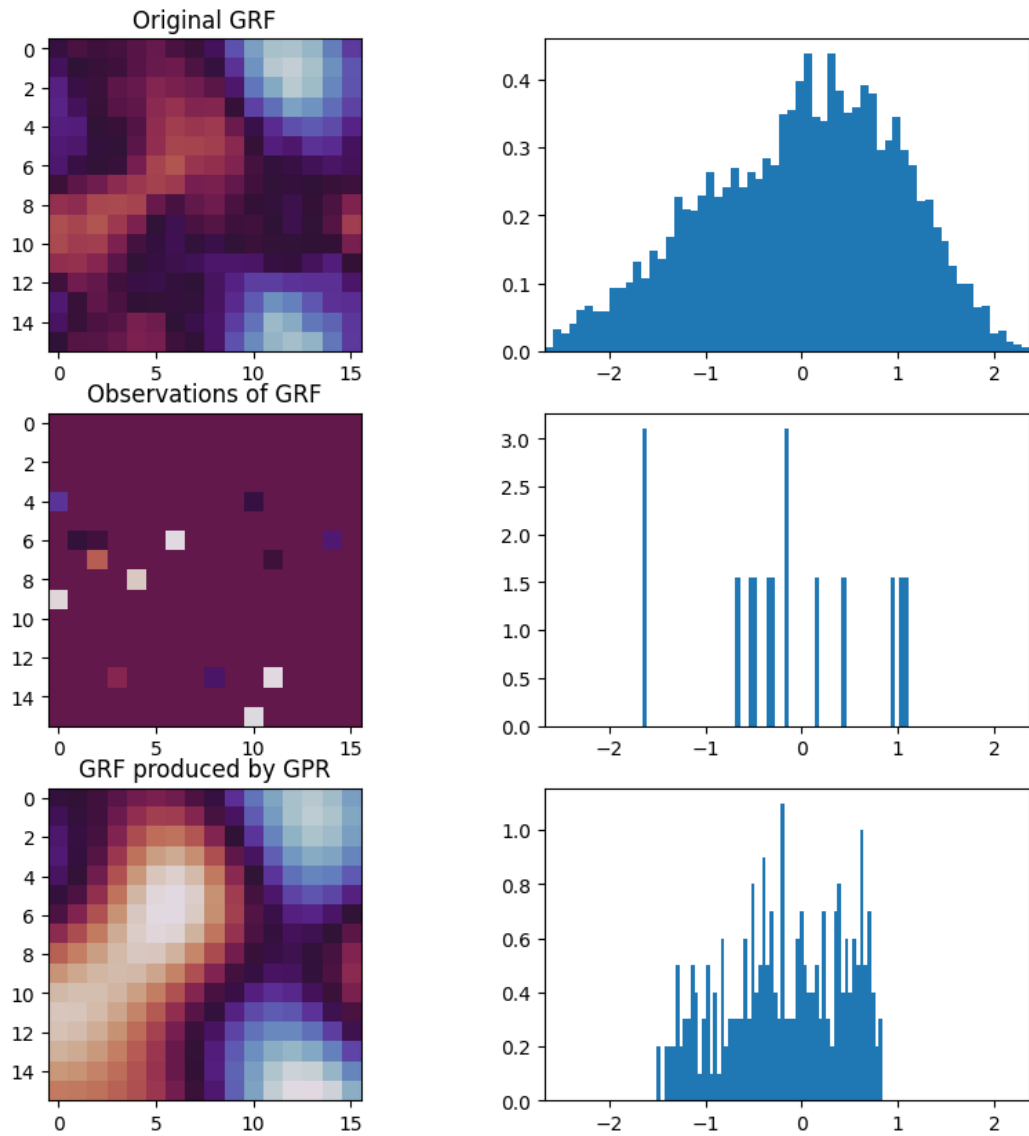
Figure 4: Gaussian process regression with $\alpha = 5.5$ and 300 observations.

Figure 5: Gaussian process regression with $\alpha = 7.5$ and 300 observations.

Figure 6: Gaussian process regression with $\alpha = 6.5$ and 20 observations.
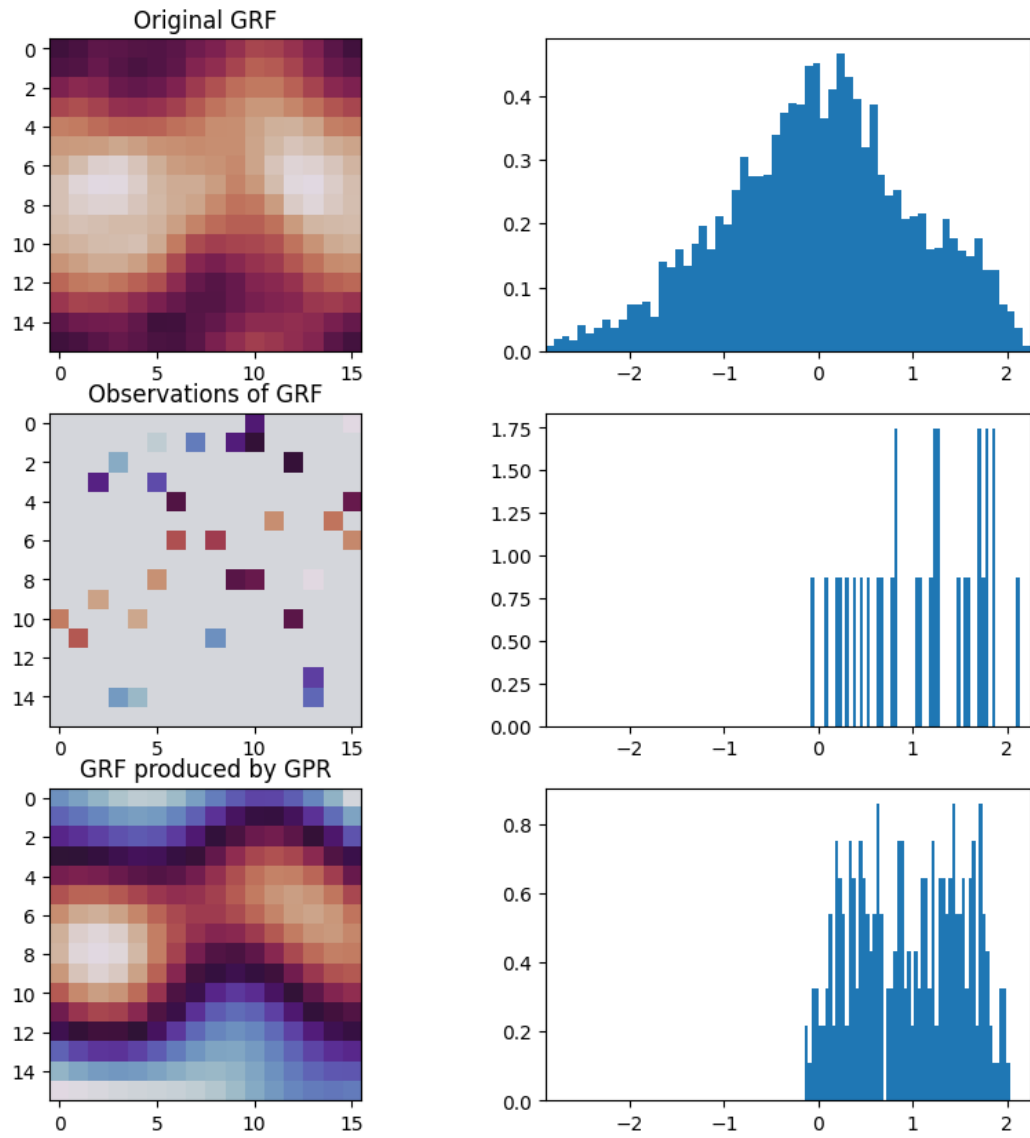
Figure 7: Gaussian process regression with $\alpha = 6.5$ and 500 observations.

Figure 8: Gaussian process regression with $\alpha = 6.5$ and 1000 observations.

References

[1] J. Görtler, R. Kehlbeck, and O. Deussen, "A Visual Exploration of Gaussian Processes," *Distill*, vol. 4, p. e17, Apr. 2019.

[2] J. Wang, "An Intuitive Tutorial to Gaussian Process Regression," *Computing in Science & Engineering*, vol. 25, pp. 4–11, July 2023. arXiv:2009.10862 [cs, stat].

[3] O. Hahn, "Cosmic Structures - Theory, Statistics, Numerics." 2024.

[4] Chappers, "Answer to "Computing Fourier transform of power law"," Mar. 2017.

[5] S.-L. , "1.7. Gaussian Processes."

[6] V. Alton, "Entonia314/CosmicStructures," July 2024. original-date: 2024-03-29T12:36:00Z.